

Completed on Tue, 07 Jun 2022

FAIL
Result

No.	Domain	Total Question	Correct	Incorrect	Unattempted
1	Cloud Concepts	14	8	6	0
2	Other	1	1	0	0
Total	All Domains	15	9	6	0

Filter By All Questions

Incorrect

A. Implement a strong identity foundation

B. Enable traceability

C. Anticipate Failure right

☐ D. Manage change in automation wrong

☐ E. Perform operations as code right

Explanation:

Answer: C, E

The operational excellence pillar of a well-architected framework has below 5 design principles.

- Perform operations as code
- Make frequent, small, reversible changes
- Refine operations procedures frequently
- Anticipate failure
- Learn from all operational failures

The security pillar of a well-architected framework has below 7 design principles.

- Implement a strong identity foundation
- Enable traceability
- Apply security at all layers
- Automate security best practices
- Protect data in transit and at rest
- Keep people away from data
- Prepare for security events

The reliability pillar of a well-architected framework has below 5 design principles.

- Automatically recover from failure
- Test recovery procedures
- Scale horizontally to increase aggregate workload availability
- Stop guessing capacity
- Manage change in automation



Option A is INCORRECT. Implement a strong identity foundation is the design principle relating to the security pillar.

Option B is INCORRECT. Enable traceability is the design principle relating to the security pillar.

Option C is CORRECT.

Option D is INCORRECT. Manage change in automation is the design principle relating to the reliability pillar.

Option E is CORRECT.

Reference:

<https://aws.amazon.com/blogs/apn/the-5-pillars-of-the-aws-well-architected-framework/>

Ask our Experts

 [View Queries](#)



Question 2

Incorrect

Domain: Cloud Concepts

What is the ability of AWS products and services to recover from disruptions and mitigate disruptions known as?

A. Resiliency right

B. Consistency

C. Durability

☐ D. Latency wrong

Explanation:

Answer: A

Resiliency is the ability to recover from disruptions and mitigate disruptions.

Consistency involves more than one system storing information, to return the same result when queried.

Durability is the system's ability to perform even upon the occurrence of unexpected events.

Latency is typically the measurement of delay between request and response.

Option A is CORRECT as Resilience is the ability of AWS products to recover from disruptions and mitigate disruptions.



Option B is INCORRECT because Consistency ensures that similar results are returned by more than one system storing information, when queried.

Option C is INCORRECT because Durability is the ability of AWS product(s) to remain functional and perform despite unexpected events' occurrence.

Option D is INCORRECT because Latency denotes the delay between getting a response after a request is made.

Reference:

<https://wa.aws.amazon.com/wat.concept.resiliency.en.html>

<https://wa.aws.amazon.com/wat.concept.consistency.en.html>

<https://wa.aws.amazon.com/wat.concept.durability.en.html>

<https://wa.aws.amazon.com/wat.concept.latency.en.html>

Ask our Experts

 View Queries



Question 3

Correct

Domain: Cloud Concepts

An architect is asked to design a solution for a distributed system in which the system's components operate in a way that components of one system do not negatively impact other components of the system.

Which of the below listed architectural best practice can help to achieve this?

- A. Request Throttling
- B. Using stateless services
- C. Enabling automatic data backup

☒ D. Implement loose coupling right

Explanation:

Answer: D

The scenario in the question talks about a distributed system wherein there is minimal-to-no dependency amongst the components. This can be achieved by implementing **loose coupling** amongst the components.



Request throttling, and, use of stateless services (Option A, B) help design resilient solution for distributed system that can withstand failures and can recover from failure quickly.

Enabling automatic data backup (Option C) is the best practice for failure management, as backups aligned with requirements will ensure the required recovery time objectives (RTO) and recovery point objectives (RPO).

Option A is INCORRECT because the scenario refers to ensure minimal dependency amongst the components. Request throttling does not ensure minimal dependency but helps build a resilient solution.

Option B is INCORRECT because using the stateless service ensures that the solution is resilient. However, loose coupling helps minimize the dependency amongst the various components in the solution.

Option C is INCORRECT because implementing automatic data backup is a failure management best practice and does not contribute to minimize component dependency.

Option D is CORRECT as loose coupling ensures that components are NOT tightly dependent on each other, assuring that if one component fails, it does not impact the working of other components.

Reference:

https://d1.awsstatic.com/whitepapers/architecture/AWS_Well-Architected_Framework.pdf Page 72, 73, 76

Ask our Experts

 [View Queries](#)



Question 4

Correct

Domain: Cloud Concepts

Which of the below statement is CORRECT regarding AWS Global infrastructure?

- ☒ A. Each AWS region has multiple Availability Zones right
- B. Many AWS regions has single availability zone
- C. Availability zones are also known as AWS Local zones
- D. To provide High Availability, AWS management console and control plane are isolated to a single region

Explanation:

Answer: A



Option A is CORRECT. The statement is Correct.

Option B is INCORRECT. The statement is incorrect. Option A is correct.

Option C is INCORRECT. The statement is incorrect. Availability zones and AWS Local Zones are different.

Option D is INCORRECT. The statement is incorrect. AWS management console and control plane utilize multi-AZs and are distributed across the AWS regions.

Reference:

<https://aws.amazon.com/about-aws/global-infrastructure/> "Regions and AZs"

Ask our Experts

 [View Queries](#)



Question 5

Correct

Domain: Cloud Concepts

An online marketplace start-up dealing in real estate is planning to move to the cloud. Which of the below is NOT a benefit of moving to the cloud?

- ☒ A. Install on a company's own servers. right
- ☐ B. Go global in minutes.
- ☐ C. Stop guessing capacity.
- ☐ D. Benefit from massive economies of scale.

Explanation:

Answer: A

Option A is CORRECT. This description belongs to on-premises and is not a benefit of moving to the cloud.

Option B is INCORRECT. This is a benefit of moving to the cloud.

Option C is INCORRECT. This is a benefit of moving to the cloud.

Option D is INCORRECT. This is a benefit of moving to the cloud.

Reference:

<https://docs.aws.amazon.com/whitepapers/latest/aws-overview/six-advantages-of-cloud-computing.html>



Ask our Experts

[+ View Queries](#)

Question 6

Incorrect

Domain: Cloud Concepts

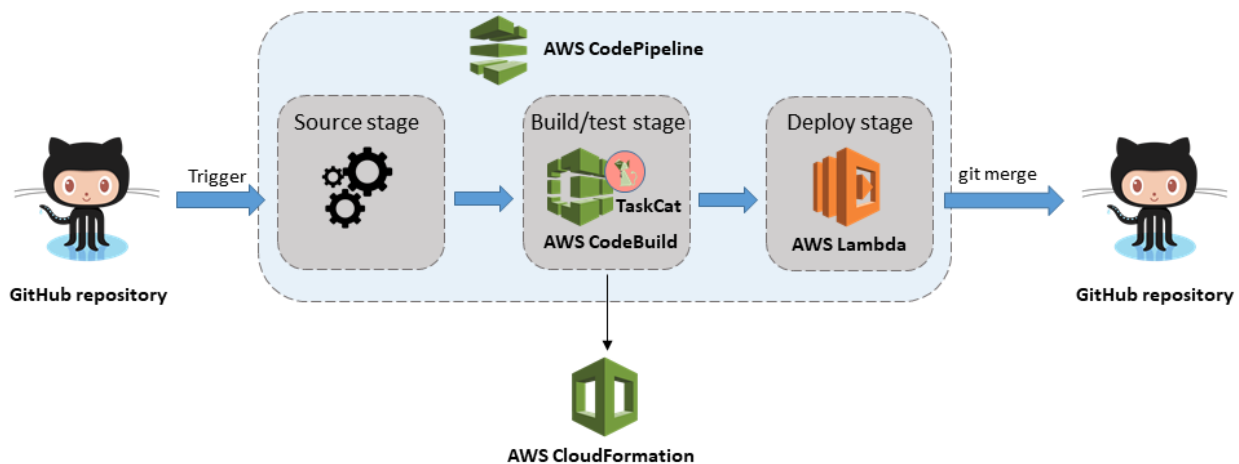
My On-Premise application's deployment cycle is roughly around 3-4 weeks. On refactoring this huge application, its features can be deployed on AWS cloud in a matter of 2-3 days. What is the benefit achieved by moving to the AWS cloud?

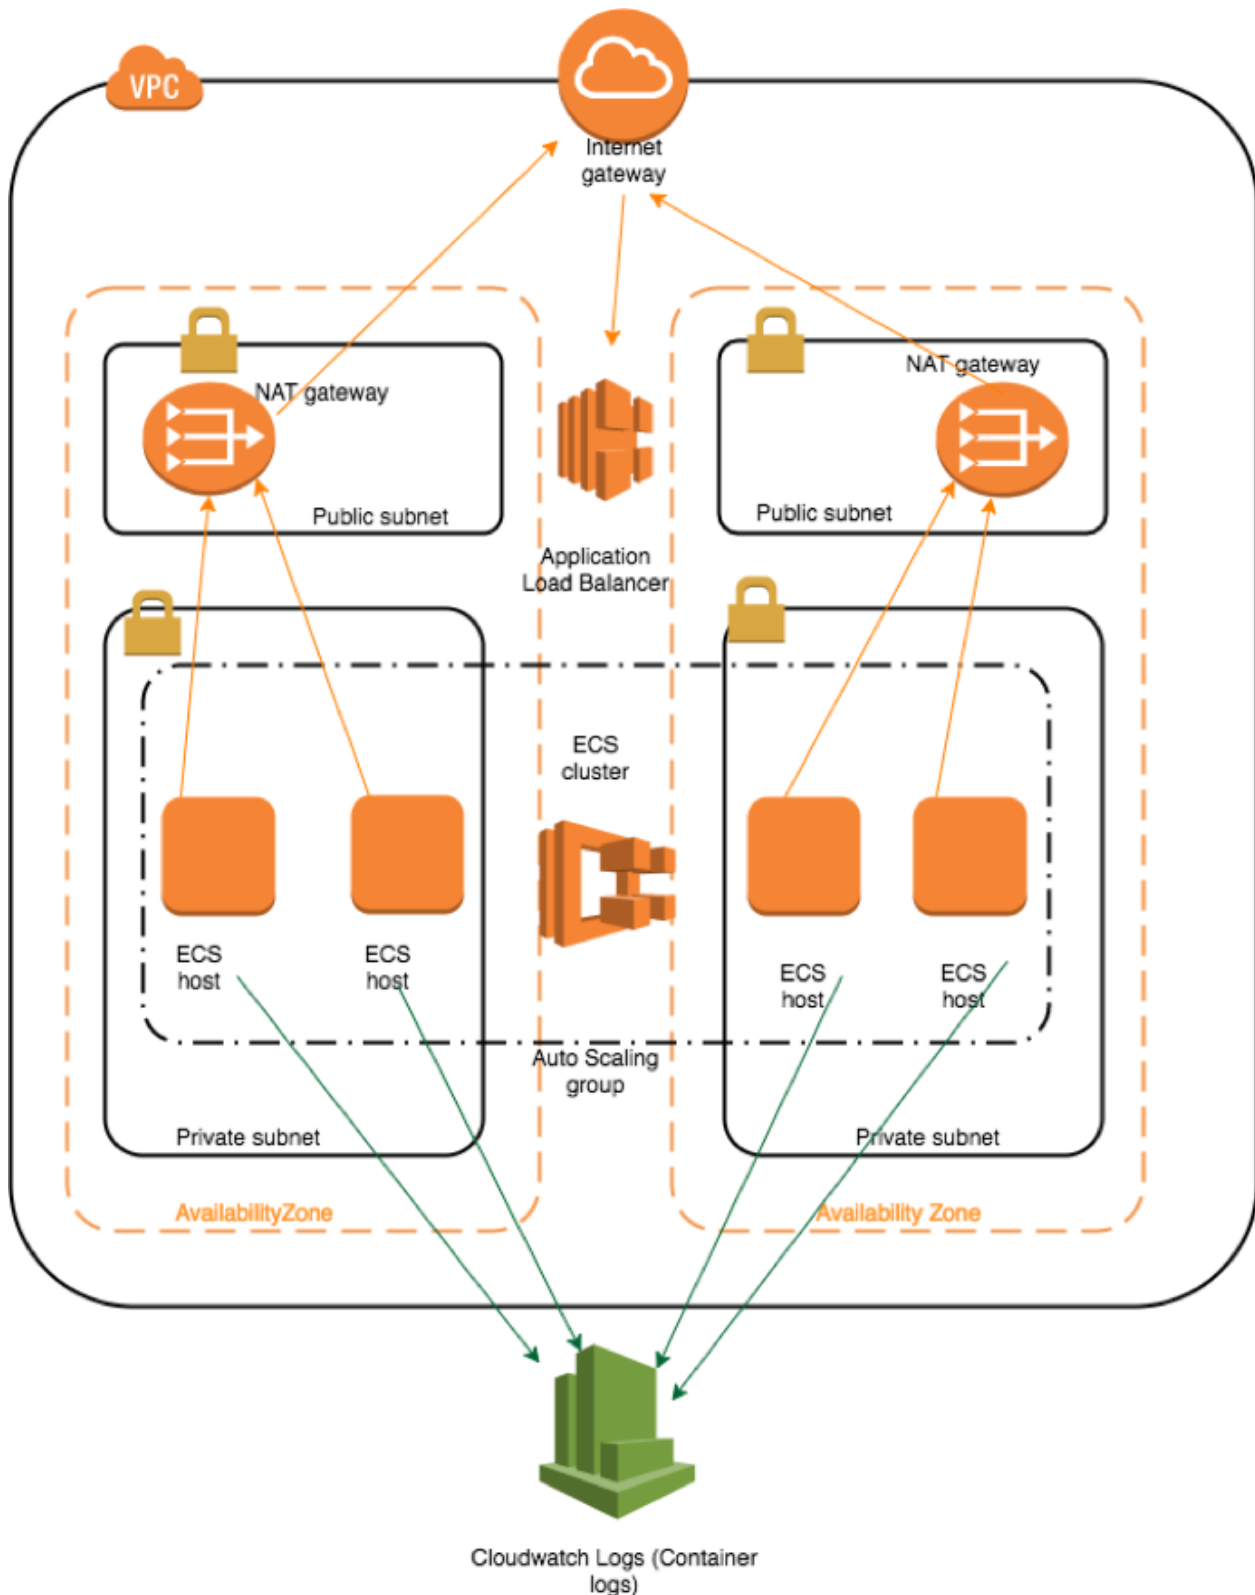
- ☐ A. Elasticity wrong
- B. Flexibility
- ☒ C. Agility right
- D. Resilience

Explanation:

Answer: C

Diagrams:





The scenario blends itself into the Business Agility value proposition where new features can be deployed faster to reduce errors. The microservices architecture paradigm assists in Business Agility, where a large Monolith application is broken down into smaller functional units that can be developed & deployed faster. AWS cloud, on the other hand, provides services like Lambda, Elastic Container Services, Elastic Kubernetes Service that assists in developing & deploying microservices & deployment tools (CI/CD) like CodeBuild, CodePipeline for automating the deployment process. Combined, they provide an environment supporting an Agile Business.

Option A is incorrect. Elasticity is the ability to scale resources on-demand whenever there is an increase in load on the existing infrastructure. This Value proposition will benefit in substantial cost savings where I need not have to guess my capacity upfront.

Option B is incorrect. Flexibility is the ability to utilize a broad range of products depending on the application & infrastructure demands. For example, there are a broad range of EC2 instance pricing models ranging from On-Demand instances to Reserved Instances to Spot Instances and EC2 instance types like General Purpose, Compute Optimized, Memory Optimized. This will help with Low or No cost entries into the AWS cloud.

Option C is CORRECT. Using an Agile environment provided by AWS Cloud, I can release application features much faster than a traditional On-Premise environment.

Option D is incorrect. Operational resiliency results in the improvement of defined SLA's and reducing unplanned outages that would result in downtime. The features of High Availability, Fault Tolerance gives rise to a resilient system.

References:

<https://www.youtube.com/watch?v=sA79G-zoSI0>

<https://www.testpreptraining.com/tutorial/aws-cloud-practitioner/define-the-aws-cloud-and-its-value-proposition/>

<https://youtu.be/FI57qGYwdK4>

Ask our Experts

 View Queries



Question 7

Correct

Domain: Cloud Concepts

I need to migrate millions of customers' financial transaction data from the On-Premise Mainframe system to a non-relational database in AWS. The database should also provide good performance for data retrieval and data analytics. Which of the following Database services is the most suitable?

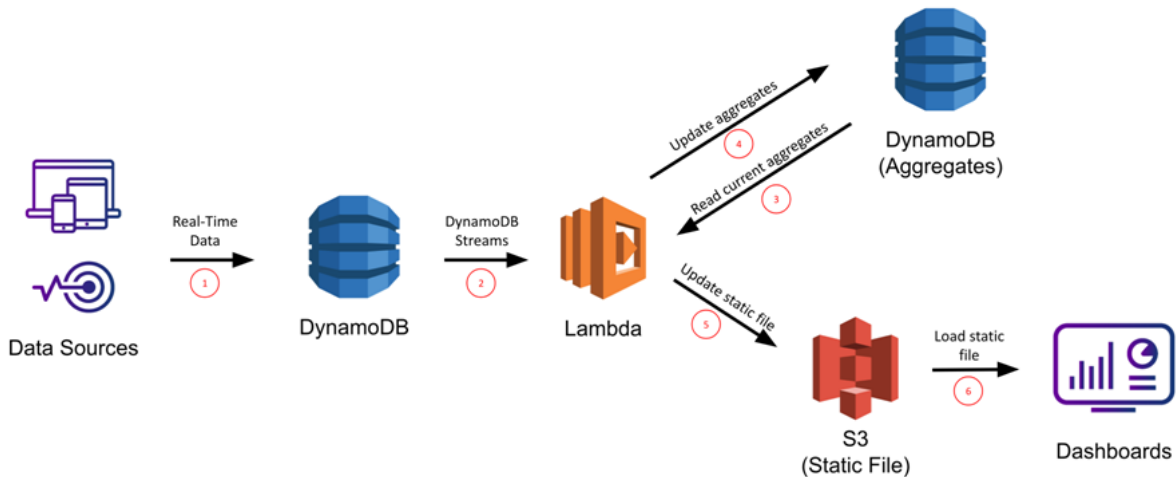
- A. Amazon RDS
- B. Amazon RedShift
- C. Amazon ElastiCache

☒ D. Amazon DynamoDB right

Explanation:

Answer: D



Diagrams:

On reading the scenario carefully, we notice that the Customer's Financial transaction data is huge. It needs storage on the cloud. NoSQL databases like DynamoDB are designed to provide seamless scalability by automatically partitioning the database as it grows in size. So a NoSQL database like DynamoDB will be the most appropriate database service that can be used for the scenario.

Option A is incorrect. Here we are exclusively talking about Huge data volumes, Data retrieval and Data analytics. RDS databases are most useful for heavy transaction processing systems. They also do not exhibit automatic partitioning capabilities with increased data volume & stream processing capabilities like a NoSQL database like DynamoDB provides.

Option B is incorrect. Amazon RedShift is a Data Warehousing solution primarily used for Operational analytics on business events. Data Warehouse may comprise a big collection of an Enterprise's structured & semi-structured data that can be used to build powerful reports & dashboards using Business Intelligence tools. Since we only have the Customer's transactional data for our scenario, RedShift will not be a good fit here.

Option C is incorrect. We are talking about the scenario for a data migration of On-Premise Mainframe data, which will require a permanent, secure data store for storing the highly sensitive Customer's financial data. Caching solutions are typically in-memory data stores used for supporting applications requiring sub-milliseconds response times. Caching solutions usually maintain a subset of the data present in a data store that does not change frequently. Also, caching solutions do not provide any facility for performing real-time data analytics, although they provide the best performance compared to any other data storage solutions.

Option D is CORRECT. DynamoDB provides DynamoDB Accelerator (DAX) which is a fully managed, highly available in-memory cache. This will help us speed up the performance of data retrieval that we require. DynamoDB also has a feature called DynamoDB streams that enables real-time capture of data changes using event notifications. This helps applications to perform analytics on real-time streaming data to build dashboards without impacting database performance. The stream events are asynchronous in nature to consuming applications like a Lambda function. Since the Customer's transactional data is highly confidential & huge in volume, a robust, scalable, secure, performant data store like DynamoDB will be the best fit for our scenario.



References:

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.Partitions.html>

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Streams.html>

Ask our Experts

 View Queries

**Question 8**

Correct

Domain: Cloud Concepts

I have certain applications On-Premise that run 24x7 & have a consistent load. I plan to move to the AWS Cloud. What is the economic feature that will benefit me most referred to?

- ☐ A. Pay-as-you-go
- ☒ B. Save when you Reserve right
- ☐ C. Pay less by using more
- ☐ D. Pay-per-compute-time

Explanation:

Answer: B

From the scenario, we can see that the On-Premises applications' workload is continuous & stable. For these kinds of applications, it is easy to predict upfront capacity. Also, since the application runs continuously, I will benefit by reserving capacity for a certain period of time.

Option A is incorrect. The Pay-as-you-go model will be best used for workloads used for short durations & have unpredictable load. On-demand EC2 instances will be the best fit for this purpose.

Option B is CORRECT. Since there is continuous usage of these applications with a predictable load, it will be best for me to reserve capacity upfront (period of 1 - 3 years) that will provide a substantial discount of 30 - 50% compared to its On Demand counterparts.

Option C is incorrect. Pay less by using more refers to volume discounts provided by AWS for increased usage. E.g., S3 Standard provides the following storage pricing, also referred to as Tiered-Pricing.

Data Storage	Storage Pricing



First TB / month	\$0.025 per GB
Next 450 TB / month	\$0.024 per GB
Next 500 TB / month	\$0.023 per GB

Option D is incorrect. Pay-per-compute-time refers to the use of serverless architectures like Lambda, where you pay only for the time when the compute resources are running. Unlike EC2 Pay-as-you-go pricing, AWS provisions resources for executing Lambda functions on the fly & removes them immediately after execution. So there is no idle utilization time that needs to be accounted for. Since our scenario consists of long-running applications, this option will be impractical for usage.

References:

<https://aws.amazon.com/pricing/#:~:text=AWS%20offers%20you%20a%20pay,utilities%20like%20water%20and%20electricity>.

<https://dzone.com/articles/the-cost-of-the-cloud-the-ultimate-aws-pricing-gui>

<https://www.apptio.com/blog/aws-reserved-instances-cost-optimization/>

https://dl.awsstatic.com/whitepapers/aws_pricing_overview.pdf

Ask our Experts

 View Queries



Question 9

Incorrect

Domain: Cloud Concepts

A client who has adopted AWS cloud services would like to ensure that his systems always scale with increasing traffic for a great end-user experience. I have implemented the same by defining AutoScaling Scale-In & Scale-Out policies & CloudWatch alarms that trigger the AutoScaling. Which Cloud Architecture Design principles have I implemented here? Select TWO most suitable options.

A. Encryption

☐ B. Operational Excellence right

C. Performance Efficiency right

☐ D. Cost Optimization wrong

E. Least privilege



Explanation:

Answers: B and C

Looking at the scenario, a good end-user experience is attached to systems being performant with increasing load on them. A combination of Load balancing & AutoScaling enables a system to handle increase in load by spinning new instances to which the load will be distributed not to saturate resources like CPU & Memory on a single server instance. For the AutoScaling itself to work efficiently, there needs to be a good monitoring system that can track resource utilization and enable automation. The scenario combines the Operational Excellence & the Performance Efficiency Architecture design principles.

Option A is incorrect. This scenario does not describe anything about data encryptions so that this option is not selected.

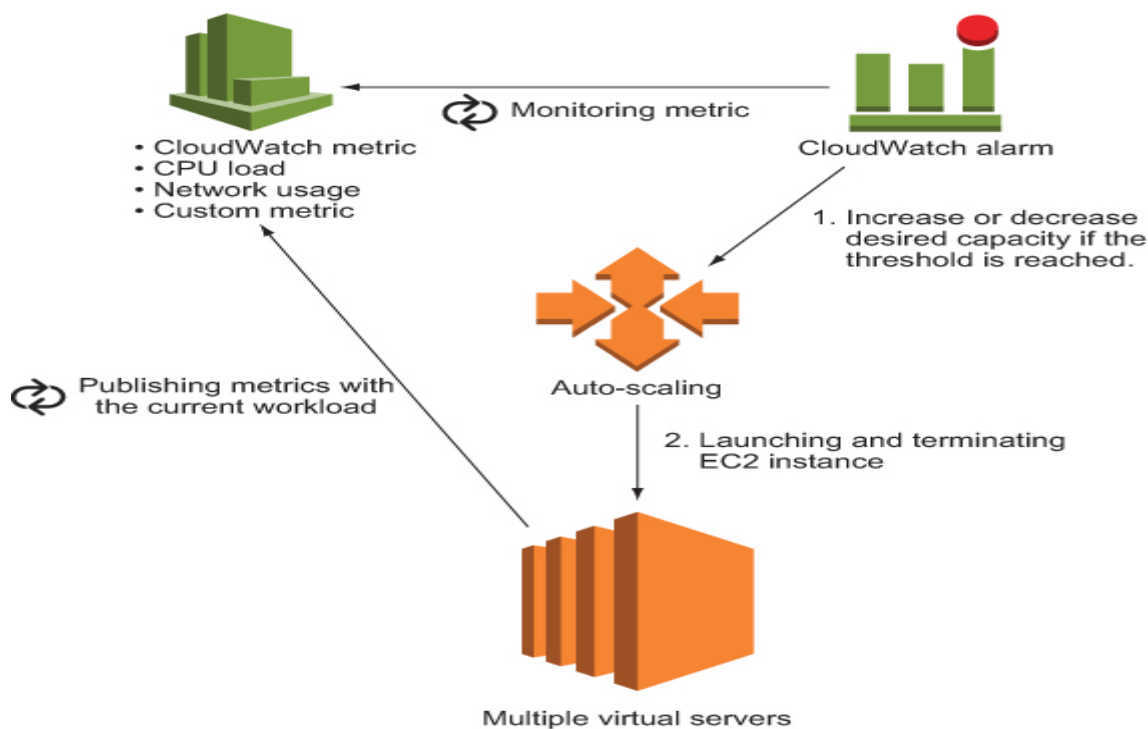
Option B is CORRECT. The monitoring mechanism used here is CloudWatch for enabling AutoScaling to happen by tracking resource utilization metrics like CPU and Memory usage. The Operational Excellence pillar focuses on Monitoring systems that will deliver continuous business value by providing automation, responding to events within the system. In our scenario, automating the system's Scalability through monitoring will help maintain the desired performance levels.

Option C is CORRECT. The Performance Efficiency pillar focuses on monitoring the performance of a system. CloudWatch metrics help monitor a system's performance by using metrics like CPU, Memory, Disk utilization etc... Automating tasks like AutoScaling using CloudWatch Alarms and defining scaling policies will ensure that the client's performance requirements will always be met with when there is an increase in traffic.

Option D is incorrect. In spite of using Scale-in & Scale-out policies, I can be benefitted from cutting down idle resource utilization costs. The scenario deals more with Performance & Operational aspects.

Option E is incorrect. Least privilege is a security principle which is not mentioned in the question.

Diagrams:

**References:**

<https://aws.amazon.com/blogs/apn/the-5-pillars-of-the-aws-well-architected-framework/>

Ask our Experts

[+ View Queries](#)

**Question 10**

Correct

Domain: Cloud Concepts

A customer's Data center and its applications are connected to AWS using a dedicated network (Direct Connect). What is the Cloud deployment model?

- ☐ A. Public Cloud
- ☐ B. Multi Cloud
- ☐ C. Private Cloud
- ☒ D. Hybrid Cloud right

Explanation:

Answer: D

Option A is incorrect. When an application is fully deployed on a vendor's (AWS, Google Cloud) Cloud Infrastructure, it is termed as a Public Cloud. The application features may have been ei



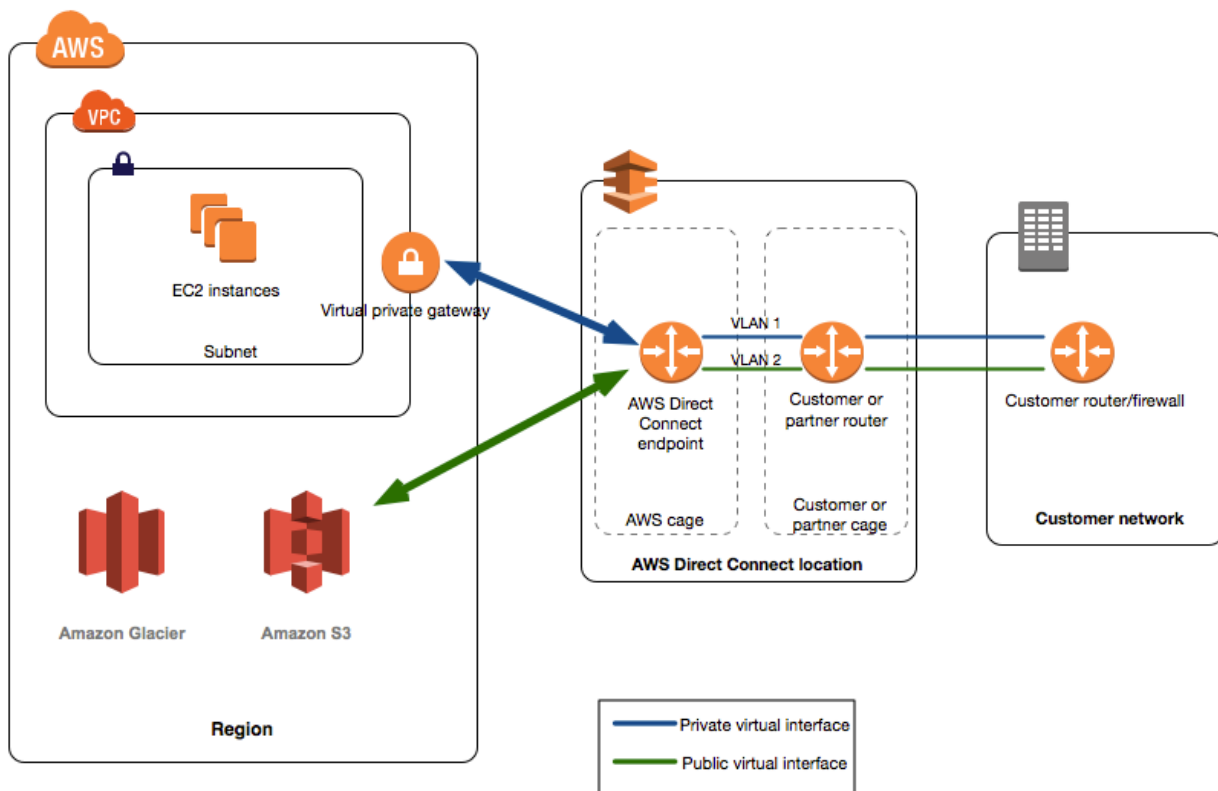
developed on the cloud (Cloud Native) or migrated from an existing Client's On-Premise infrastructure. Our scenario uses both the Customer's Data Center and AWS.

Option B is incorrect. A multi-cloud leverages different Cloud Provider's environments. He may deploy his applications on AWS and GoogleCloud-based on compute speed requirements, availability of managed services etc...With the use of Multi-Cloud, it is best to use Open Source tools like Terraform for Infrastructure as Code, Splunk for monitoring that are cloud vendor-agnostic to avoid vendor lockin. Since our scenario uses only the Customer's Data Center and AWS, there are no multiple cloud deployments.

Option C is incorrect. A Private Cloud refers to a cloud environment that has been built In house on the Client's premises using Virtualization and resource management tools. Private clouds are usually built for specific security and regulatory requirements, specially hardware or configuration requirements, and extremely critical network latency that may not be possible on a vendor's public cloud platform. Since our scenario has both the Client's data center & the AWS Cloud, it will not result in a Private Cloud configuration.

Option D is CORRECT. The cloud deployment model is a Hybrid Cloud. Hybrid clouds are often used where there is a requirement of burst capacity where workloads are "spilled over" to a different cloud environment to meet capacity demands for a short period of time. Here, purchasing capacity may not be a good idea since it will be extremely ineffective from a cost standpoint resulting in under utilization of resources once the requirement for capacity ends. The other compelling use case for a Hybrid cloud is a Highly Available / Disaster Recovery environment where the cloud model offers a lot of flexibility in making resources in the event of a Datacenter failure where the client may not need upfront investments for resources for an alternate site.

Diagrams:



References:

<https://youtu.be/XlqoKLL040A>https://youtu.be/X_ZZO-oeNeghttps://youtu.be/9KZL0_NuiUU[Ask our Experts](#)[+ View Queries](#)

Question 11

Correct

Domain: Other

0

Your company has set up EC2 Instances in a VPC for their application. The IT Security department needs to understand what security mechanisms are available to protect the Instances when it comes to traffic going in and out of the instance. What are the two layers of security provided by AWS in the VPC? (Select TWO)

☐ A. Security Groups right☐ B. Network ACLs right

C. DHCP Options

D. Route Tables

Explanation:

Correct Answers: A and B

The AWS Documentation mentions the following

A *security group* acts as a virtual firewall for your instance to control inbound and outbound traffic. When you launch an instance in a VPC, you can assign up to five security groups to the instance.

A *network access control list (ACL)* is an optional layer of security for your VPC that acts as a firewall for controlling traffic in and out of one or more subnets. You might set up network ACLs with rules similar to your security groups in order to add an additional layer of security to your VPC.

Option C is incorrect since this is used to decide on the DNS servers for the VPC

Option D is incorrect since this is used for routing traffic in the VPC

For more information on VPC security groups and Network ACL's, please visit the below URL

https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_SecurityGroups.htmlhttps://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_ACLS.html

Ask our Experts

[+ View Queries](#)

Question 12

Correct

Domain: Cloud Concepts

Which of the following may NOT be an Economic benefit to a client using AWS cloud services?

- ☒ A. The Client is running a dedicated MySQL Database Server on AWS with his own CPU bound license (BOYL). right
- B. The Client is running Spot Instances for batch data processing workloads.
- C. The client is running applications with a relatively predictable & consistent resource Demand using AWS Reserved Instances.
- D. The client is using S3 Intelligent Tiering storage class while uploading objects.
- E. The client is using an Active - Passive failover routing strategy of his On - Premise Data Center to AWS cloud.

Explanation:

Answers: A

Option A is CORRECT. CPU bound software licenses will require a Dedicated Host tenancy model rather than a Shared tenancy model for the EC2 instances hosting the MySQL Database software. Dedicated hosts are the most expensive tenancy model when it comes to pricing. For example, an m4 large On Demand dedicated host is 24 times more expensive than an On-demand shared host. A dedicated host tenancy may be used under exceptional circumstances of spiky traffic, while purchasing Hardware On-Premise may not be the best option. Then the client can take advantage of Cloud Elasticity.

Option B is incorrect. Batch processing jobs rely more on accuracy rather than on speed which forms a good use case for using Spot Instances, providing economies of compute where intermittent disconnections may not be a real problem.

Option C is incorrect. AWS Reserved Instances have discounts on the EC2 usage. This method can provide an economic benefits.

Option D is incorrect. S3 Intelligent Tiering is a smart solution for managing the lifecycle of S3 objects resulting in economies of cost. With intelligent tiring, there are no retrieval fees nor there are any fees for moving objects between tiers unlike using S3 Lifecycle policies which incurs data transfer charges. S3 Lifecycle policies are also often challenging to define due to the unpredictable nature of application adoption and usage. Even in scenarios where access



frequency is known, it may so happen that customers may not use proper storage class adjustments resulting in non-optimized budgets.

Option E is incorrect. An Active-Passive failover will always be economical to a client using the strategy for Disaster Recovery scenarios since he will not be investing in an entire redundant site with all resources running simultaneously. He will have the flexibility to select his strategy depending on his Recovery Time objectives. This will help him to save on costs.

References:

<https://www.liquidweb.com/blog/private-cloud/>

<https://www.nutanix.com/blog/s3-intelligent-tiering-aws-cost-saving>

<https://www.xplenty.com/blog/the-5-types-of-data-processing/#batch>

<https://aws.amazon.com/ec2/dedicated-hosts/pricing/>

<https://aws.amazon.com/ec2/pricing/on-demand/>

Ask our Experts

 View Queries



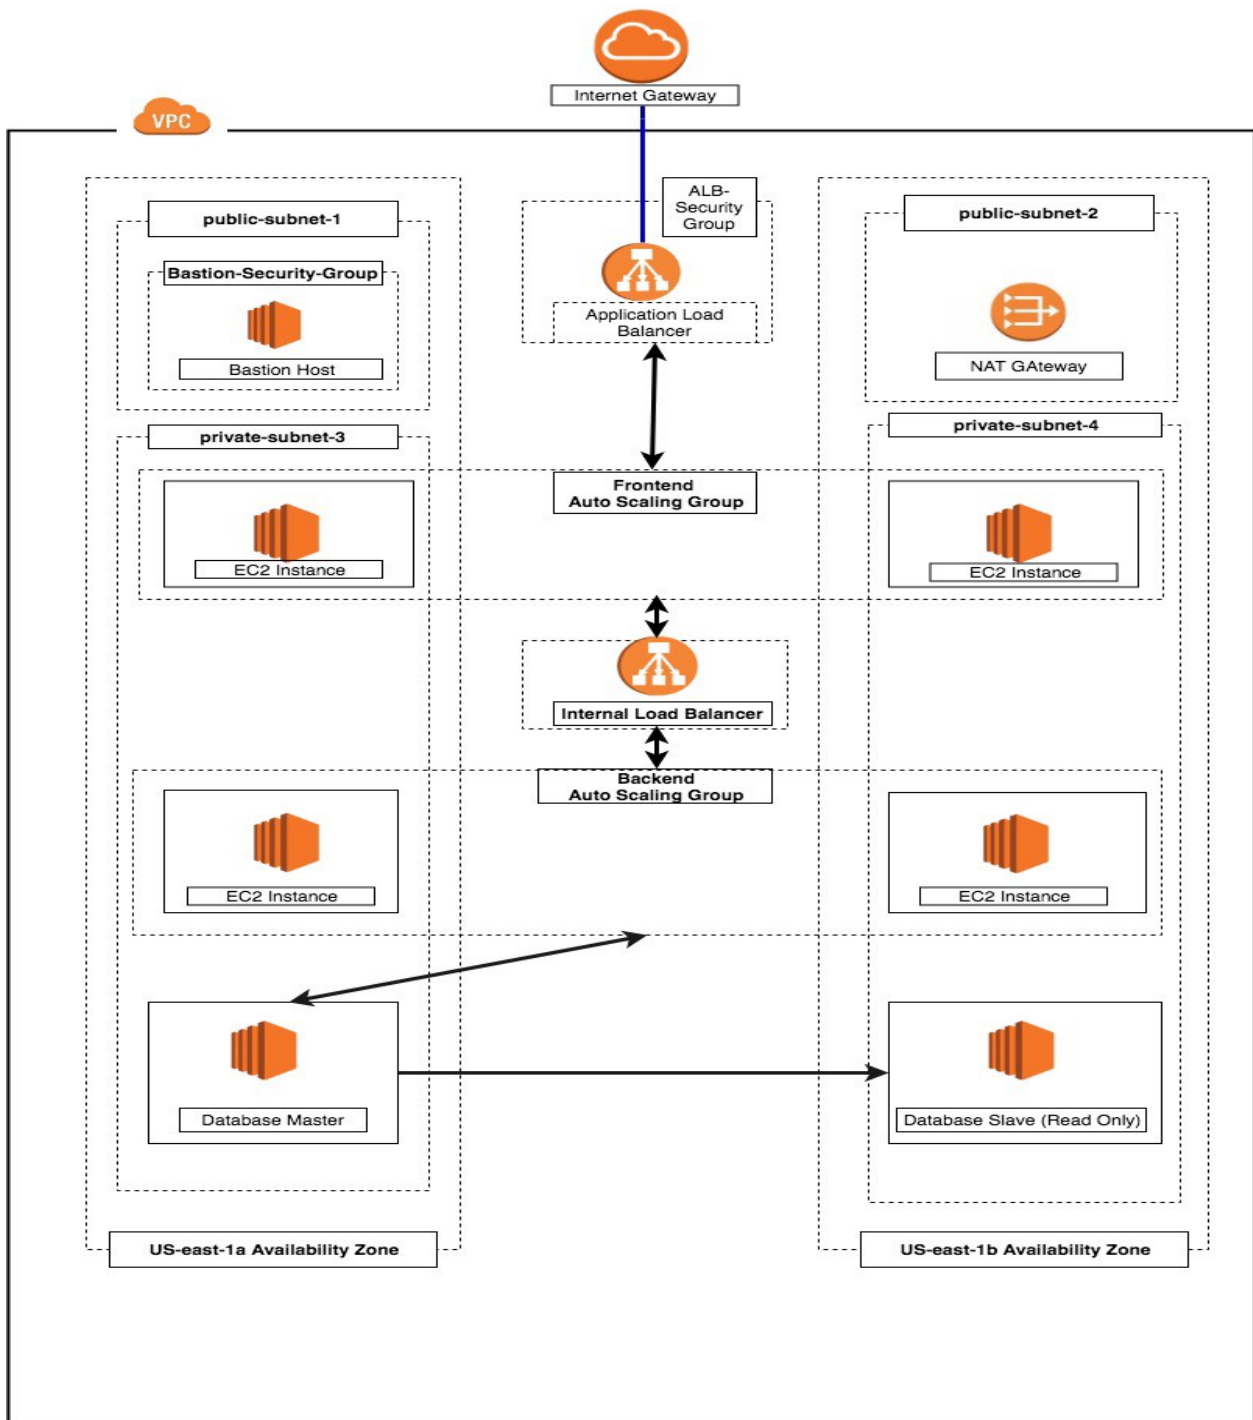
Question 13

Incorrect

Domain: Cloud Concepts

Which of the following indicates a Single Point Of Failure (SPOF) in an AWS Cloud Architecture design?
Refer to the figure below.





A. Application Load Balancer

B. Bastion Host right

C. EC2 Instance

☐ D. NAT Gateway within an Availability Zone wrong

Explanation:

Answer: B

Option A is incorrect. For an Application Load Balancer, you need to specify at least two subnets each in a different AZ. This ensures that the load balancer will redirect load to a different AZ if one of them goes down ensuring HA and redundancy. Application Load balancer performs Health

Checks on its registered targets and will route requests only to Healthy Targets. Even though the ALB is depicted as a single resource in the figure, multiple ALB instances are created for each AZ behind the scenes. It ensures that the ALB is not a SPOF.

Option B is CORRECT. We can see in the figure that the Bastion Host has been deployed on only one of the AZ's causing a SPOF. For redundancy, the Bastion Host needs to be created in Multiple AZ's where a failure in a single AZ will not impact access to Server instances in Private subnets through the Bastion Hosts.

Option C is incorrect. We can see EC2 instances being deployed on Private Subnets in multiple AZ's, ensuring redundancy & avoiding SPOF.

Option D is incorrect. NAT Gateway offers redundancy within a single AZ. Hence it is not a SPOF within the AZ.

References:

<https://packetswitch.co.uk/aws-nat-gateway-high-availability/#:~:text=NAT%20Gateway%20HA%20scenario,Availability%20Zones%20lose%20Internet%20access.>

<https://docs.aws.amazon.com/vpc/latest/userguide/vpc-nat-gateway.html>

<https://jayendrapatil.com/aws-bastion-host/>

<https://stackoverflow.com/questions/46698011/are-amazon-elastic-load-balancer-elb-failure-proof>

Ask our Experts

 View Queries



Question 14

Correct

Domain: Cloud Concepts

0

Which of the following AWS resources or the AWS features (cloud concepts) does NOT provide automation capabilities?

A. AWS Elastic Beanstalk

B. Amazon DynamoDB

C. AWS CloudFormation

☒ D. RDS manual snapshot right

Explanation:



Answer: D

Option A is incorrect. Elastic Beanstalk provides a fast way to deploy a web application on AWS. It automatically handles resource provisioning, load balancing, autoscaling, and monitoring when configured.

Option B is incorrect. DynamoDB contains a feature called DynamoDB streams that provides change events for automatically capturing data during operations like CREATE, DELETE, UPDATE on its tables.

Option C is incorrect. AWS CloudFormation has automation capabilities, is used to automate workload deployments on the AWS Cloud.

Option D is CORRECT. Unlike automated backups, manual snapshots are taken care of by users when needed. It is not an automation method.

References:

<https://docs.aws.amazon.com/prescriptive-guidance/latest/migration-operations-integration/aws-services-for-automation.html>

Ask our Experts

 [View Queries](#)

**Question 15**

Incorrect

Domain: Cloud Concepts

I have certain applications On-Premises that experience times within a year where infrastructure takes a heavier load impact (e.g., Christmas, Thanksgiving, etc.) than other times in the year. You do not want to decommission the on-premises infrastructure. What is the easiest and most cost-effective way in which I can handle this load?

- ☐ A. By moving all my infrastructure to AWS Cloud and using On-Demand capacity
- ☐ B. By creating a Private Cloud environment in my On-Premises data center that will provide me with the required elasticity
- ☒ C. By using Scheduled Reserved Instances to match capacity reservation for the load wrong
- ☐ D. By provisioning Burst Capacity on the AWS Cloud for the duration of the load right

Explanation:

Answer: D

Option A is incorrect. On looking at the scenario, we see that the variable load is only for a specific duration of time. So moving the infrastructure entirely to a Public Cloud will not be the best



solution to gain maximum benefit out of the elastic nature of a Public Cloud.

Option B is incorrect. A Private Cloud will be more beneficial where there is a consistent load rather than a variable load. It will be good to have a Private Cloud hosting the applications for economies of cost and agility rather than elasticity, which can be best obtained using a Public Cloud.

Option C is incorrect. Reserved Instances are usually best chosen where there is consistent usage and predictable load for a certain duration of time (1 - 3 Years). Scheduled Reserved Instances have the ability to reserve capacity for a predictable recurring schedule that may be in a day, week, or month. The advantage here is w.r.t the costs that I incur for Reserved Instances that can be managed without paying everything upfront rather than elasticity.

Option D is CORRECT. This is the best way to reduce the costs of purchasing Hardware and getting the benefits of the elasticity and On-Demand pricing provided by a Public Cloud environment. On these specific occasions, I can demand a burst capacity by going to the AWS Public cloud. That will certainly help me maintain the performance of my applications at heavy load. Once the load reduces, I can then terminate the instances that are no longer used to save costs.

References:

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/burstable-performance-instances.html>

https://youtu.be/OzVpynd_2GM

Ask our Experts

 [View Queries](#)



Finish Review

Certification

Cloud Certification

Java Certification

PM Certification

Big Data Certification

Support

Contact Us

Help Topics

Company

Become Our Instructor

Support

Discussions

Blog

Business



Join us on Slack!

Join our open **Slack community** and get your queries answered instantly! Our experts are online to answer your questions!



